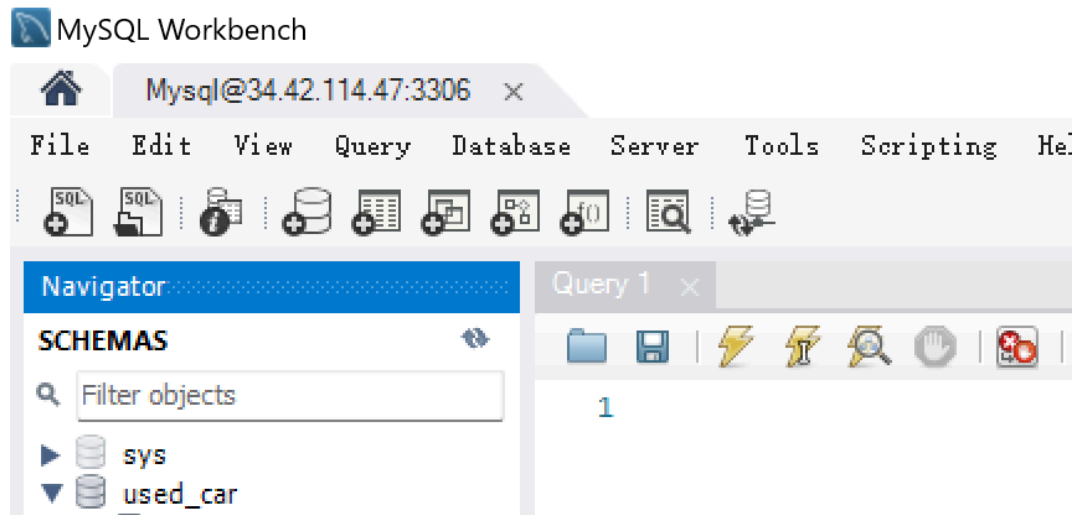


# Stage 3: Database Implementation and Indexing

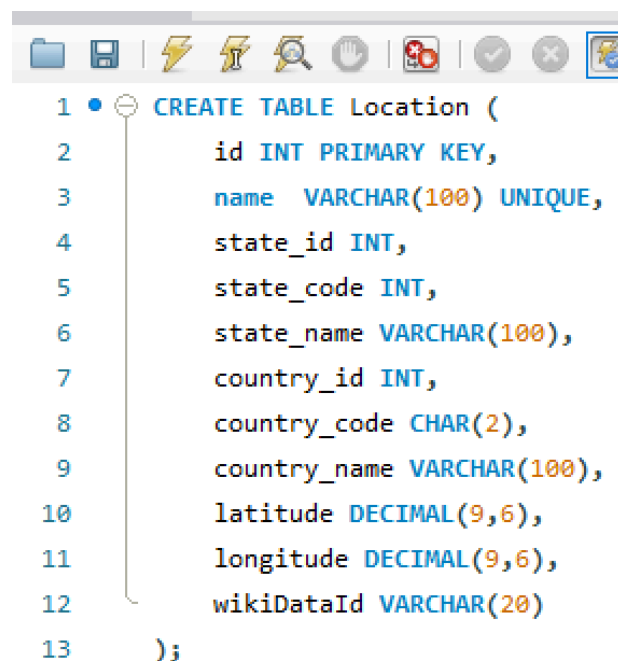
## I. Database implementation

access to the database(from MySQL Workbench)



DDL commands

- Table Location



100% 19:4

Result Grid Filter Rows: Search Export:

Field	Type	Null	Key	Default	Extra
id	int	NO	PRI	NULL	
name	varchar(100)	YES	UNI	NULL	
state_id	int	YES		NULL	
state_code	int	YES		NULL	
state_name	varchar(100)	YES		NULL	
country_id	int	YES		NULL	
country_code	char(2)	YES		NULL	
country_name	varchar(100)	YES		NULL	
latitude	decimal(9,6)	YES		NULL	
longitude	decimal(9,6)	YES		NULL	
wikiDataId	varchar(20)	YES		NULL	

- Table User

```
CREATE TABLE User (
  id int PRIMARY KEY,
  name VARCHAR(30),
  email VARCHAR(40),
  birthDate DATETIME,
  password VARCHAR(50),
  locations VARCHAR(50),
  phone VARCHAR(30),
  FOREIGN KEY (locations) REFERENCES Location(name)
);
```

100% 19:4

Result Grid Filter Rows: Search Export:

Field	Type	Null	Key	Default	Extra
id	int	NO	PRI	NULL	
name	varchar(100)	YES	UNI	NULL	
state_id	int	YES		NULL	
state_code	int	YES		NULL	
state_name	varchar(100)	YES		NULL	
country_id	int	YES		NULL	
country_code	char(2)	YES		NULL	
country_name	varchar(100)	YES		NULL	
latitude	decimal(9,6)	YES		NULL	
longitude	decimal(9,6)	YES		NULL	
wikiDataId	varchar(20)	YES		NULL	

- Table Brand

Query 1

```

1 CREATE TABLE Brand (
2     Maker VARCHAR(30) PRIMARY KEY,
3     Founder VARCHAR(30),
4     CompanyType VARCHAR(30),
5     Headquarters VARCHAR(30),
6     FoundedYear INT
7 );

```

Result Grid

Filter Rows:

Export:

Wrap Cell

	Field	Type	Null	Key	Default	Extra
▶	Maker	varchar(30)	NO	PRI	NULL	
	Founder	varchar(30)	YES		NULL	
	CompanyType	varchar(30)	YES		NULL	
	Headquarters	varchar(30)	YES		NULL	
	Producta	varchar(50)	YES		NULL	
	FoundedYear	int	YES		NULL	
	Industry	varchar(50)	YES		NULL	
	ProductionOutput	varchar(50)	YES		NULL	
	Revenue	int	YES		NULL	
	OperatingIncome	int	YES		NULL	
	NetIncome	int	YES		NULL	
	TotalAsset	int	YES		NULL	
	TotalEquity	int	YES		NULL	
	NumberOfEmplo...	int	YES		NULL	

- Table Car

```

1  CREATE TABLE Car (
2      CarId INT AUTO_INCREMENT PRIMARY KEY,
3      Make VARCHAR(300),
4      Model VARCHAR(300),
5      CarTitle VARCHAR(200),
6      CarSubTitle VARCHAR(500),
7      CarPrice INT,
8      ManufactureYear INT,
9      BodyType VARCHAR(300),
10     Mileage INT,
11     EngineVolume DECIMAL(10, 2),
12     EngineSize DECIMAL(10, 2),
13     TransmissionType VARCHAR(300),
14     FuelType VARCHAR(200),
15     TotalPreviousOwners INT,
16     FOREIGN KEY (Make) REFERENCES Brand(Maker)
17 );
  
```

Result Grid | Filter Rows: | Export: | Wrap Cell Content:

	Field	Type	Null	Key	Default	Extra
▶	CarId	int	NO	PRI	NULL	auto_increment
	Make	varchar(300)	YES	MUL	NULL	
	Model	varchar(300)	YES		NULL	
	CarTitle	varchar(200)	YES	UNI	NULL	
	CarSubTitle	varchar(500)	YES		NULL	
	CarPrice	int	YES		NULL	
	ManufactureYear	int	YES		NULL	
	BodyType	varchar(300)	YES		NULL	
	Mileage	int	YES		NULL	
	EngineVolume	decimal(10,2)	YES		NULL	
	EngineSize	decimal(10,2)	YES		NULL	
	TransmissionType	varchar(300)	YES		NULL	
	FuelType	varchar(200)	YES		NULL	
	TotalPreviousO...	int	YES		NULL	

- Table Advertisement

Limit to 1000 rows

```

1 CREATE TABLE Advertisement (
2   AdvertisementId INT AUTO_INCREMENT PRIMARY KEY,
3   CarId INT,
4   UserId INT,
5   CarAttentionGrabber VARCHAR(300),
6   FinanceAvailable BOOLEAN,
7   Discounted BOOLEAN,
8   FOREIGN KEY (CarId) REFERENCES Car(CarId),
9   FOREIGN KEY (UserId) REFERENCES User(UserId)
10 );

```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

	Field	Type	Null	Key	Default	Extra
▶	AdvertisementId	int	NO	PRI	<b>NULL</b>	auto_increment
	CarId	int	YES	MUL	<b>NULL</b>	
	UserId	int	YES	MUL	<b>NULL</b>	
	CarAttentionGrabber	varchar(300)	YES		<b>NULL</b>	
	FinanceAvailable	tinyint(1)	YES		<b>NULL</b>	
	Discounted	tinyint(1)	YES		<b>NULL</b>	

- Table Rating

Limit to 1000 rows

Open a script file in this editor

```

1 CREATE TABLE Rating (
2   RatingId INT AUTO_INCREMENT PRIMARY KEY,
3   CarTitle VARCHAR(100),
4   Price VARCHAR(50),
5   OverallRating DECIMAL(2,1),
6   ExteriorRating DECIMAL(2,1),
7   InteriorRating DECIMAL(2,1),
8   RideQuality DECIMAL(2,1),
9   Foreign Key (CarTitle) references Car(CarTitle)
10 );

```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

	Field	Type	Null	Key	Default	Extra
▶	RatingId	int	NO	PRI	<b>NULL</b>	auto_increment
	CarTitle	varchar(100)	YES	MUL	<b>NULL</b>	
	Price	varchar(50)	YES		<b>NULL</b>	
	OverallRating	decimal(2,1)	YES		<b>NULL</b>	
	ExteriorRating	decimal(2,1)	YES		<b>NULL</b>	
	InteriorRating	decimal(2,1)	YES		<b>NULL</b>	
	RideQuality	decimal(2,1)	YES	YES	<b>NULL</b>	

Data volume

```
1 • select count(*) from Location;
```

Result Grid		Filter Rows:	Export:
	count(*)		
▶	128905		

1 • select count(\*) from User;

Result Grid		Filter Rows:
	count(*)	
▶	1000	

1 • select count(\*) from Brand;

Result Grid		Filter Rows:
	count(*)	
▶	132	

1 • select count(\*) from Car;

Result Grid		Filter Rows:
	count(*)	
▶	3669	

 Limit

```
1 • select count(*) from Advertisement;
```

Result Grid		Filter Rows:	Export
	count(*)		
▶	30538		

```
1 • select count(*) from Rating;
```

Result Grid		Filter Rows:
	count(*)	
▶	56	

## II. Advanced Queries

### Query 1: Average Reviews of Different Brands

- **Techniques used:** JOIN, GROUP BY
- **Purpose:** Calculate the average overall rating for each car brand based on user reviews.

```
mysql> SELECT c.Make, AVG(r.OverallRating) AS AvgOverallRating
-> FROM Car c
-> JOIN Rating r ON c.CarTitle = r.CarTitle
-> GROUP BY c.Make
-> ORDER BY AvgOverallRating DESC
-> LIMIT 15;
```

Make	AvgOverallRating
Volkswagen	5.00000
Jaguar	5.00000
Audi	4.45000
Lamborghini	4.10000
MINI	4.00000
Isuzu	3.80000
MG	3.70000
Hyundai	2.61250
Jeep	2.46667
Porsche	2.00000
Toyota	1.66667
SKODA	1.33333
Land Rover	0.71429
BMW	0.62500
Aston Martin	0.00000

15 rows in set (0.03 sec)

### Query 2: Discounted cars priced above average

- **Techniques used:** JOIN, Subquery
- **Purpose:** Identify cars with discounts that are still priced above the average price of all cars.

```
mysql> SELECT a.AdvertisementId, c.CarTitle, c.CarPrice
-> FROM Advertisement a
-> JOIN Car c ON a.CarId = c.CarId
-> WHERE a.Discounted = 1
-> AND c.CarPrice > (
-> SELECT AVG(CarPrice) FROM Car
-> )
-> ORDER BY c.CarPrice DESC
-> LIMIT 15;
```

AdvertisementId	CarTitle	CarPrice
4224	BMW M8 Competition Convertible	148555
4226	BMW M8 Competition Convertible	148555
4242	BMW M8 Gran Coupe	147070
4233	BMW M8 Competition Gran Coupe	141465
4234	BMW M8 Competition Gran Coupe	141465
4235	BMW M8 Competition Gran Coupe	141465
4237	BMW M8 Competition Gran Coupe	141465
4238	BMW M8 Competition Gran Coupe	141465
4822	BMW X6 M Competition	136855
4825	BMW X6 M Competition	136855
4228	BMW M8 Competition Coupe	133297
4229	BMW M8 Competition Coupe	133297
4230	BMW M8 Competition Coupe	133297
4231	BMW M8 Competition Coupe	133297
4232	BMW M8 Competition Coupe	133297

15 rows in set (0.03 sec)

### Query 3: Find the average price at which each brand is listed in ads

- **Techniques used:** JOIN, GROUP BY
- **Purpose:** Compute the average listing price for each brand based on all ads.



```
mysql> SELECT c.Make, AVG(c.CarPrice) AS AvgListedPrice
-> FROM Car c
-> JOIN Advertisement a ON c.CarId = a.CarId
-> GROUP BY c.Make
-> ORDER BY AvgListedPrice DESC
-> LIMIT 15;
```

Make	AvgListedPrice
Bugatti	895000.0000
Lamborghini	251352.0462
Ferrari	236756.1200
McLaren	217940.0000
Aston Martin	173017.1207
Rolls-Royce	139725.7895
De Tomaso	124995.0000
Porsche	121404.6435
Radical	120000.0000
Bentley	111094.9545
Bac	99990.0000
Bowler	98491.0000
Crendon	94748.0000
Factory Five	89950.0000
Shelby	79995.0000

15 rows in set (0.07 sec)

#### Query4: High-price SUVs above brand average

- **Techniques used:** Subquery, filtering conditions
- **Purpose:** List SUV models with engine volume > 1.4L that are priced above their brand's average.

```
mysql> SELECT c.CarTitle, c.Make, c.CarPrice
-> FROM Car c
-> WHERE c.CarPrice > (
->     SELECT AVG(c2.CarPrice)
->     FROM Car c2
->     WHERE c2.Make = c.Make
-> ) AND EngineVolume > 1.4 AND BodyType = 'suv'
-> ORDER BY c.CarPrice DESC
-> LIMIT 15;
```

CarTitle	Make	CarPrice
Rolls-Royce Cullinan Black Badge	Rolls-Royce	379950
Rolls-Royce Cullinan	Rolls-Royce	299000
Mercedes-Benz Maybach Gls Class	Mercedes-Benz	249987
Mercedes-Benz G Series	Mercedes-Benz	229996
Bentley Bentayga S	Bentley	204890
Bentley Bentayga	Bentley	194995
Ford F150	Ford	187140
Mercedes-Benz G63 AMG	Mercedes-Benz	187000
Bentley Bentayga Hybrid	Bentley	182950
Bentley Bentayga Mulliner	Bentley	175890
Land Rover New Range Rover	Land Rover	159891
BMW ALPINA XB7	BMW	157768
BMW X7	BMW	150665
Land Rover Range Rover Efi	Land Rover	150000
Mercedes-Benz G Class Amg Station Wagon	Mercedes-Benz	149506

15 rows in set (0.16 sec)

#### Query 5: Overall Ratings of Car Models Priced Above the Average

- **Techniques used:** JOIN, Subquery, ORDER BY

- **Purpose:** Evaluate whether expensive cars have better overall ratings.

```
mysql> SELECT c.CarTitle, c.CarPrice, r.OverallRating
-> FROM Car c
-> JOIN Rating r ON c.CarTitle = r.CarTitle
-> WHERE c.CarPrice > (
->     SELECT AVG(CarPrice) FROM Car
-> )
-> ORDER BY OverallRating DESC, EngineSize
-> LIMIT 15;
```

CarTitle	CarPrice	OverallRating
Land Rover Range Rover Evoque	52500	5.0
Lamborghini Huracan Evo	229000	4.1
Porsche 911	72000	4.0
Porsche Cayenne	76648	4.0
Jeep Compass	40099	3.7
Jeep Compass	40099	3.7
Hyundai IONIQ 5	39985	0.0
Land Rover Range Rover Velar	39600	0.0
Land Rover Range Rover Velar	39600	0.0
Land Rover Defender	69995	0.0
Land Rover Defender	69995	0.0
Lexus ES	39995	0.0
Volvo XC60	39705	0.0
Porsche Macan	63090	0.0
Jeep Wrangler	59752	0.0

15 rows in set (0.01 sec)

#### Query 6: Low-Mileage Cars with Fewer Owners, Ordered by Earliest Manufacture Year

- **Techniques used:** Subquery, GROUP BY, ORDER BY

- **Purpose:** Find used cars with mileage below average and fewer than 3 previous owners, ordered by the earliest manufacture year.

```
mysql> SELECT CarTitle
-> FROM Car
-> WHERE Mileage < (
->     SELECT AVG(Mileage)
->     FROM Car
-> )
-> AND TotalPreviousOwners < 3
-> GROUP BY CarTitle
-> ORDER BY MIN(ManufactureYear)
-> LIMIT 15;
```

```
+-----+
| CarTitle |
+-----+
| Audi A5 Coupe |
| Audi RS Q3 |
| BMW 7 Series Saloon (LWB) |
| Audi S3 Saloon |
| Audi Tts Coupe |
| BMW 2 Series Gran Coupe |
| BMW 7 Series Saloon |
| Abarth 695C |
| Abarth Abarth 500 |
| Alpine Alpine |
| BMW 4 Series Coupe |
| Audi TT Roadster |
| Audi S8 |
| BMW 8 Series Coupe |
| BMW ALPINA B8 |
+-----+
15 rows in set (0.13 sec)
```

### III. Indexing Analysis

#### A. Query 1

Screenshot for Analyzing the Query

```


18 • EXPLAIN ANALYZE
19 SELECT c.Make, AVG(r.OverallRating) AS AvgOverallRating
20 FROM Car c
21 JOIN Rating r ON c.CarTitle = r.CarTitle
22 GROUP BY c.Make
23 ORDER BY AvgOverallRating DESC
24 LIMIT 15;

```

130%	1:16	Result Grid	Filter Rows: Search	Export:
EXPLAIN				
-> Limit: 15 row(s) (actual time=3.19..3.2 rows=15 loops=1) -> Sort: AvgOverallRating DESC, limit input to 15 row(s) per ch				

## Indexing Strategies Explored

We experimented with indexing different attributes that appeared in the **JOIN**, **WHERE**, **GROUP BY**, or **HAVING** clauses. Below are the results for each configuration:



## Performance Observations

Surprisingly, all indexing configurations yielded the **same query cost of 25.5**, indicating **no observable performance gain or degradation** regardless of which attribute was indexed.

### Why indexing didn't help:

**Small Dataset Size:** The Rating table has only 56 data, which is too small to perform cost variation with different indexing. As a result, the dataset is small enough that the query planner determines a full table scan is more efficient than using indexes.

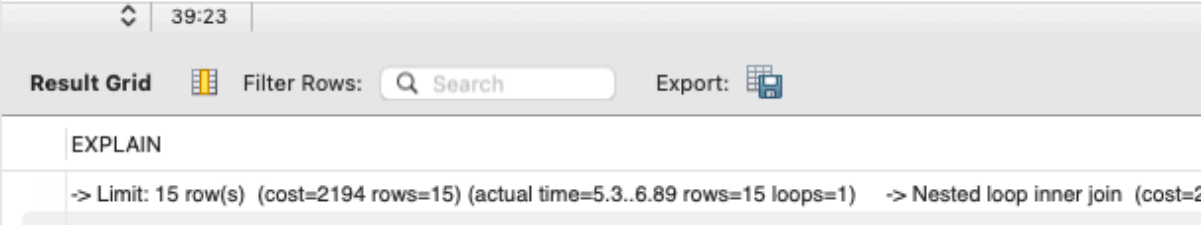
## Final Index Design Decision

Given that no index yielded performance improvement, **no additional indexes are selected** for this specific query. We opt to avoid unnecessary indexing to save on write/update overhead and storage cost.

### B. Query 2

#### Screenshot for Analyzing the Query

```
29 • EXPLAIN ANALYZE
30 SELECT a.AdvertisementId, c.CarTitle, c.CarPrice
31 FROM Advertisement a
32 JOIN Car c ON a.CarId = c.CarId
33 WHERE a.Discounted = 1
34 AND c.CarPrice > (
35     SELECT AVG(CarPrice) FROM Car
36 )
37 LIMIT 15;
38
```




The screenshot shows a database query editor with a SQL query and its execution plan. The query is an `EXPLAIN ANALYZE` statement for a JOIN query. The execution plan shows a nested loop inner join with a cost of 2194 and 15 rows returned.

EXPLAIN
-> Limit: 15 row(s) (cost=2194 rows=15) (actual time=5.3..6.89 rows=15 loops=1) -> Nested loop inner join (cost=2

## Indexing Strategies Explored

We experimented with indexing different attributes that appeared in the `JOIN`, `WHERE`, `GROUP BY`, or `HAVING` clauses. Below are the results for each configuration:



## Performance Observations

The results clearly show that **adding indexes significantly reduces the query cost**.

- Indexing **Advertisement(CarId)** alone reduces the cost to **1439**, as it optimizes the **JOIN**.
- Indexing **Advertisement(Discounted)** reduces it further to **695**, helping the **WHERE** clause filter on the discounted flag.
- Indexing **Car(CarPrice)** improves filtering for price-based logic in both the **WHERE** clause and the subquery.
- The best result was achieved using a **composite index on Advertisement(Discounted, CarId)**, dropping the cost to **577** — likely due to MySQL taking advantage of multi-column filtering and join optimization.

## Final Index Design Decision

We selected the **composite index on Advertisement(Discounted, CarId)** as the optimal index for this query.


It provides the lowest cost by efficiently supporting:

- Filtering by **Discounted**
- The **JOIN** with **Car** on **CarId**

### C. Query 3


#### Screenshot for Analyzing the Query

```
37  EXPLAIN ANALYZE
38  SELECT c.Make, AVG(c.CarPrice) AS AvgPrice
39  FROM Advertisement a
40  JOIN Car c ON a.CarId = c.CarId
41  GROUP BY c.Make
42  HAVING AvgPrice > (
43      SELECT AVG(CarPrice) FROM Car
44  )
45  ORDER BY AvgPrice DESC
46  LIMIT 15;
47
```



#### Indexing Strategies Explored

We experimented with indexing different attributes that appeared in the **JOIN**, **WHERE**, **GROUP BY**, or **HAVING** clauses. Below are the results for each configuration:



## Performance Observations

Overall, **none of the indexing strategies significantly improved the query performance**. The slight drop from 5111 to 5004 with the `Car(CarPrice)` index was the best observed improvement, but it's too minor to justify the index alone.

- Indexing `Car(CarPrice)` helped slightly, likely due to its use in the subquery and `AVG()` computation.
- Composite indexes such as `Car(Make, CarPrice)` actually **worsened** performance to **6114**, suggesting the planner did not find them helpful, or they added overhead.
- Other composite indexes such as `Car(CarId, Make)` or `Car(CarId, CarPrice)` made **no difference** compared to individual indexes.

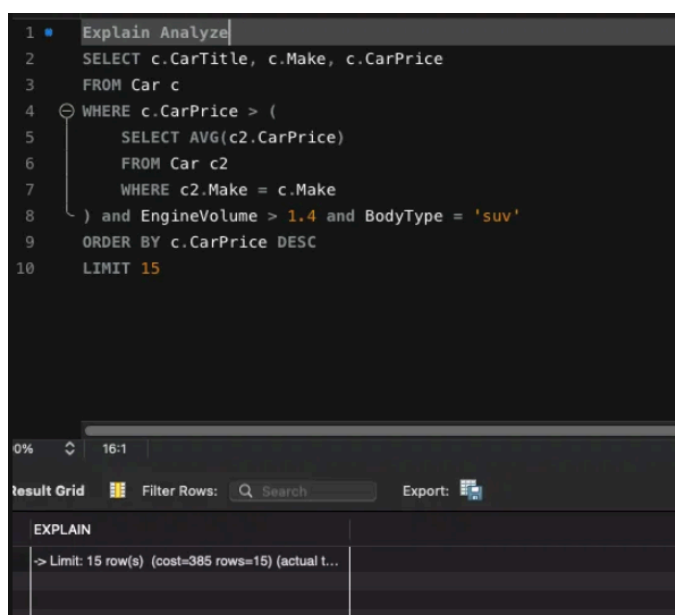
## Final Index Design Decision

Since **none of the indexes produced meaningful performance improvements**, we recommend **not adding any new index** for this query.

We will retain the default (or possibly `Car(CarPrice)`) only if needed elsewhere in the workload.

### D. Query 4

#### ScreenShot for Analyzing the Query



```
1 Explain Analyze
2 SELECT c.CarTitle, c.Make, c.CarPrice
3 FROM Car c
4 WHERE c.CarPrice > (
5     SELECT AVG(c2.CarPrice)
6     FROM Car c2
7     WHERE c2.Make = c.Make
8 ) and EngineVolume > 1.4 and BodyType = 'suv'
9 ORDER BY c.CarPrice DESC
10 LIMIT 15
```

The screenshot shows a database query editor with a query and its execution plan. The query is a complex one involving a subquery and multiple filters. The execution plan is shown in a table below the query.

EXPLAIN
-> Limit: 15 row(s) (cost=385 rows=15) (actual t...



## Performance Observations


### 1. Before Indexing

- Top-level Cost: **385**
- Scan Type: Full Table Scan
- Filter Estimated Rows: **3606**

### 2. Indexing Experiments

Index	Top-level Cost	Notes
On <b>CarPrice</b> DESC ( <b>idx_price</b> )	1.55	Dramatic cost reduction. Used index scan.
On <b>EngineVolume</b>	385	No improvement, full table scan remained.
On <b>BodyType</b> ( <b>idx_BodyType</b> )	107	Partial improvement. Reduced scanned rows to 1038.

### 3. Analysis



- Indexing **CarPrice** provided the most dramatic improvement.

- BodyType indexing helped, but less significantly.
- EngineVolume index had no effect, likely due to low selectivity.

## Final Index Design Decision

**Selected: Car(CarPrice DESC)**

- Reason: Reduced cost from 385 to 1.55 by leveraging index for filtering and sorting.

## E. Query 5

### ScreenShot for Analyzing the Query

```

1 Explain Analyze
2 SELECT c.CarTitle, c.CarPrice, r.OverallRating
3 FROM Car c
4 JOIN Rating r ON c.CarTitle = r.CarTitle
5 WHERE c.CarPrice > (
6     SELECT AVG(CarPrice)
7     FROM Car
8 )
9 ORDER BY OverallRating DESC, EngineSize
10 LIMIT 15;
  
```

The screenshot shows a database query analyzer interface. The top pane displays the SQL query for 'Query 5'. The bottom pane shows the 'EXPLAIN' output, indicating that the query is limited to 15 rows and took 2.29 seconds to execute.

## Performance Observations

### 1. Before Indexing

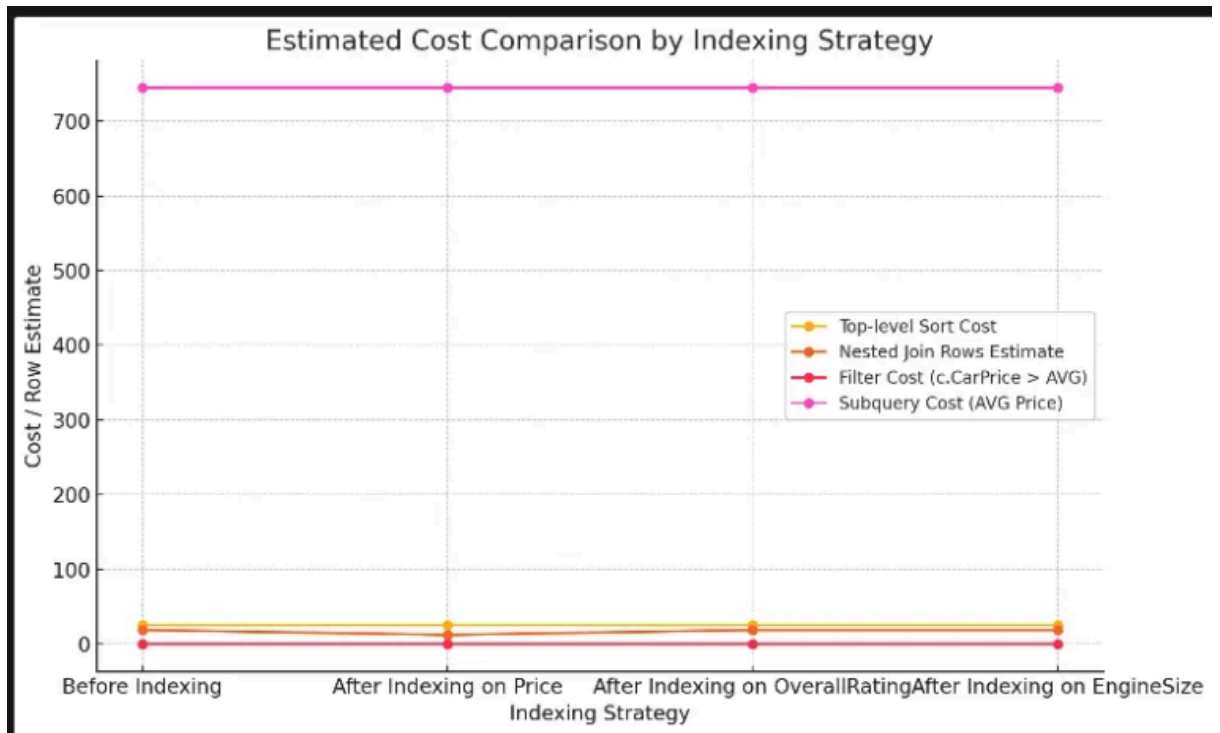
- Top-level Cost: 25.5
- Subquery Cost: 745
- Table Scan on Car: 385

### 2. Indexing Experiments

Index	Top-level Cost	Notes
On <b>CarPrice</b>	25.5	Slight improvement in filter rows

		(from 0.333 → 0.218)
On <b>Rating.CarTitle</b>	25.5	No improvement, Rating table is small.
On <b>EngineSize</b>	25.5	No impact. Not used for filtering.

### 3. Analysis



- The index on **CarPrice** improved filter selectivity slightly.
- Other indexes had no visible impact due to small table size or no filtering role.

### Final Index Design Decision

- Selected: **Car(CarPrice)**
- Reason: Slight improvement in join cost and filter accuracy.  
Others ineffective

## F. Query 6

### ScreenShot for Analyzing the Query

```
3 Explain analyze
4 SELECT CarTitle
5 FROM Car
6 WHERE Mileage < (
7     SELECT AVG(Mileage) FROM Car
8 )
9 AND TotalPreviousOwners < 3
10 GROUP BY CarTitle
11 ORDER BY MIN(ManufactureYear)
12 LIMIT 15;
```

100% 10:12

Result Grid Filter Rows: Search

EXPLAIN
-> Limit: 15 row(s) (actual time=13.7..13.7 rows...

### Observations

#### 1. Before Indexing


- Top-level Cost: 345
- Estimated Rows: 401

#### 2. Indexing Experiments

Index	Top-level Cost	Filter Rows	Notes
On Mileage	460	749	Cost increased
On TotalPreviousOwners	282	1049	Best cost reduction, despite high row count.
On Manufacture	345	201	No change. Used in ORDER BY but no

Year			filter.
------	--	--	---------

## Analysis



- Index on **TotalPreviousOwners** reduced cost most.
- Index on **Mileage** worsened performance (overhead).
- ManufactureYear index unused for optimization.

## Final Index Design Decision

- Selected: **Car(TotalPreviousOwners)**
- Reason: Lowered top-level cost from 345 → 282. Best tradeoff overall.